

AI Planning

Introduction to Artificial Intelligence CSCE476/876, Spring 2009:

Reading

- Required reading
 - Sections 11.1—11.4
- Recommended reading
 - AIMA Section 10.3: Actions, Situations, and Events
 - Chapter 11 entirely

Outline

- Background
 - Situation Calculus
 - Frame, qualification, & ramification problems
- Representation language
- Algorithms

Background

- Focus
 - The focus here is deterministic planning
 - Environment is fully observable
 - Results of actions is deterministic
 - Relaxing the above requires dealing with uncertainty
 - Problem types: sensorless, contingency, exploration
- Planning ‘communities’ in AI
 - Logic-based: Reasoning About Actions & Change
 - Less formal representations: Classical AI Planning
 - Uncertainty (UAI): Graphical Models such as
 - Markov Decision Processes (MDP), Partially Observable MDPs, etc.
- AI Planning is **not** MRP (Material Requirements Planning)


Actions, events, and change

- Planning requires a representation of time
 - to express & reason about sequences of actions
 - to express the effects of actions on the world
- Propositional Logic
 - does not offer a representation for time
 - Each action description needs to be repeated for each step
- Situation Calculus (AIMA Section 10.3)
 - Is based on FOL
 - Each time step is a ‘situation’
 - Allows to represent plans and reason about actions & change

Situation Calculus: Ontology


- Situations
- Fluents
- Atemporal (or eternal) predicates & functions

AIMA Section 10.3



Situation Calculus: Ontology

- Situations
 - Initial state: S_0
 - A function $Result(a.s)$ gives the situation resulting from applying action a in situation s
- Fluents
 - Functions & predicates whose truth values can change from one situation to the other
 - Example: $\neg Holding(G_1, S_0)$
- Atemporal (or eternal) predicates and functions
 - Example: $Gold(G_1)$, $LeftLegOf(Wumpus)$



Situation Calculus

- Sequence of actions
 - $\text{Result}([], s) = s$
 - $\text{Result}([a|\text{seq}], s) = \text{Result}(\text{seq}, \text{Result}(a, s))$
- Projection task
 - Deducing the outcome of a sequence of actions
- Planning task
 - Find a sequence of actions that achieves a desired effect

Example: Wumpus World

- Fluents
 - $\text{At}(o,p,s), \text{Holding}(o,s)$
- Agent is in [1,1], gold is in [1,2]
 - $\text{At}(\text{Agent},[1,1],S_0) \wedge \text{At}(G_1,[1,2],S_0)$
- In S_0 , we also need to have:
 - $\text{At}(o,x,S_0) \Leftrightarrow [(o=\text{Agent}) \wedge x=[1,1]] \vee [(o=G_1) \wedge x=[1,2]]$
 - $\neg \text{Holding}(o,S_0)$
 - $\text{Gold}(G_1) \wedge \text{Adjacent}([1,1],[1,2]) \wedge \text{Adjacent}([1,2],[1,1])$
- The query is:
 - $\exists \text{seq } \text{At}(G_1,[1,1],\text{Result}(\text{seq},S_0))$
- The answer is
 - $\text{At}(G_1,[1,1],\text{Result}(\text{Go}([1,1],[1,2]),\text{Grab}(G_1),\text{Go}([1,2],[1,1]),S_0))$

Importance of Situation Calculus

- Historical note
 - Situation Calculus was the first attempt to formalizing planning in FOL
 - Other formalisms include Event Calculus
 - The area of using logic for planning is informally called in the literature “Reasoning About Action & Change”
- Highlighted three important problems
 1. Frame problem
 2. Qualification problem
 3. Ramification problem

'Famous' Problems

- Frame problem
 - Representing all things that stay the same from one situation to the next
 - Inferential and representational
- Qualification problem
 - Defining the circumstances under which an action is guaranteed to work
 - Example: what if the gold is slippery or nailed down, etc.
- Ramification problem
 - Proliferation of implicit consequences of actions as actions may have secondary consequences
 - Examples: How about the dust on the gold?

Outline

- Background
 - Situation Calculus
 - Frame, qualification, & ramification problems
- **Representation language**
- Algorithms

Planning Languages

- Languages must represent..
 - States
 - Goals
 - Actions
- Languages must be
 - Expressive for ease of representation
 - Flexible for manipulation by algorithms

State Representation

- A state is represented with a conjunction of positive literals
- Using
 - Logical Propositions: $Poor \wedge Unknown$
 - FOL literals: $At(Plane1, OMA) \wedge At(Plane2, JFK)$
- FOL literals must be ground & function-free
 - **Not allowed**: $At(x, y)$ or $At(Father(Fred), Sydney)$
- Closed World Assumption
 - What is not stated are assumed false

Goal Representation

- Goal is a partially specified state
- A proposition satisfies a goal if it contains all the atoms of the goal and possibly others..
 - Example: $\text{Rich} \wedge \text{Famous} \wedge \text{Miserable}$ satisfies the goal $\text{Rich} \wedge \text{Famous}$

Action Representation

- Action Schema

- Action name
- Preconditions
- Effects

At(WHI, LNK), Plane(WHI),
Airport(LNK), Airport(OHA)

Fly(WHI, LNK, OHA)

At(WHI, OHA), \neg At(WHI, LNK)

- Example

Action(Fly(p, from, to),

PRECOND: $\text{At}(p, \text{from}) \wedge \text{Plane}(p) \wedge \text{Airport}(\text{from}) \wedge \text{Airport}(\text{to})$

EFFECT: $\neg \text{At}(p, \text{from}) \wedge \text{At}(p, \text{to})$

- Sometimes, Effects are split into ADD list and DELETE list

Applying an Action

- Find a substitution list θ for the variables
 - of all the precondition literals
 - with (a subset of) the literals in the current state description
- Apply the substitution to the propositions in the effect list
- Add the result to the current state description to generate the new state
- Example:
 - Current state: $\text{At}(\text{P1}, \text{JFK}) \wedge \text{At}(\text{P2}, \text{SFO}) \wedge \text{Plane}(\text{P1}) \wedge \text{Plane}(\text{P2}) \wedge \text{Airport}(\text{JFK}) \wedge \text{Airport}(\text{SFO})$
 - It satisfies the precondition with $\theta = \{p/\text{P1}, \text{from}/\text{JFK}, \text{to}/\text{SFO}\}$
 - Thus the action $\text{Fly}(\text{P1}, \text{JFK}, \text{SFO})$ is applicable
 - The new current state is: $\text{At}(\text{P1}, \text{SFO}) \wedge \text{At}(\text{P2}, \text{SFO}) \wedge \text{Plane}(\text{P1}) \wedge \text{Plane}(\text{P2}) \wedge \text{Airport}(\text{JFK}) \wedge \text{Airport}(\text{SFO})$

Languages for Planning Problems

- STRIPS
 - Stanford Research Institute Problem Solver
 - Historically important
- ADL
 - Action Description Languages
 - See Table 11.1 for STRIPS versus ADL
- PDDL
 - Planning Domain Definition Language
 - Revised & enhanced for the needs of the International Planning Competition
 - Currently [version 3.1](#)

Example: Air Cargo

- See Figure 11.2
- Initial state
- Goal State
- Actions: Load, Unload, Fly

Example: Spare Tire Problem

- See Figure 11.3
- Initial State
- Goal State
- Actions:
 - *Remove(Spare, Trunk), Remove(Flat, Axle)*
 - *PutOn(Spare, Axle)*
 - *LeaveOvernight*
- Note
 - the negated precondition $\neg At(Flat, Axle)$ not allowed in STRIPS.
 - Could be easily replaced with *Clear(Axle)*, adding one more predicate to the language

Example: Blocks World

- See Fig 11.4
- Initial state
- Goal
- Actions:
 - *Move(b,x,y)*
 - *MoveToTable(b,x)*

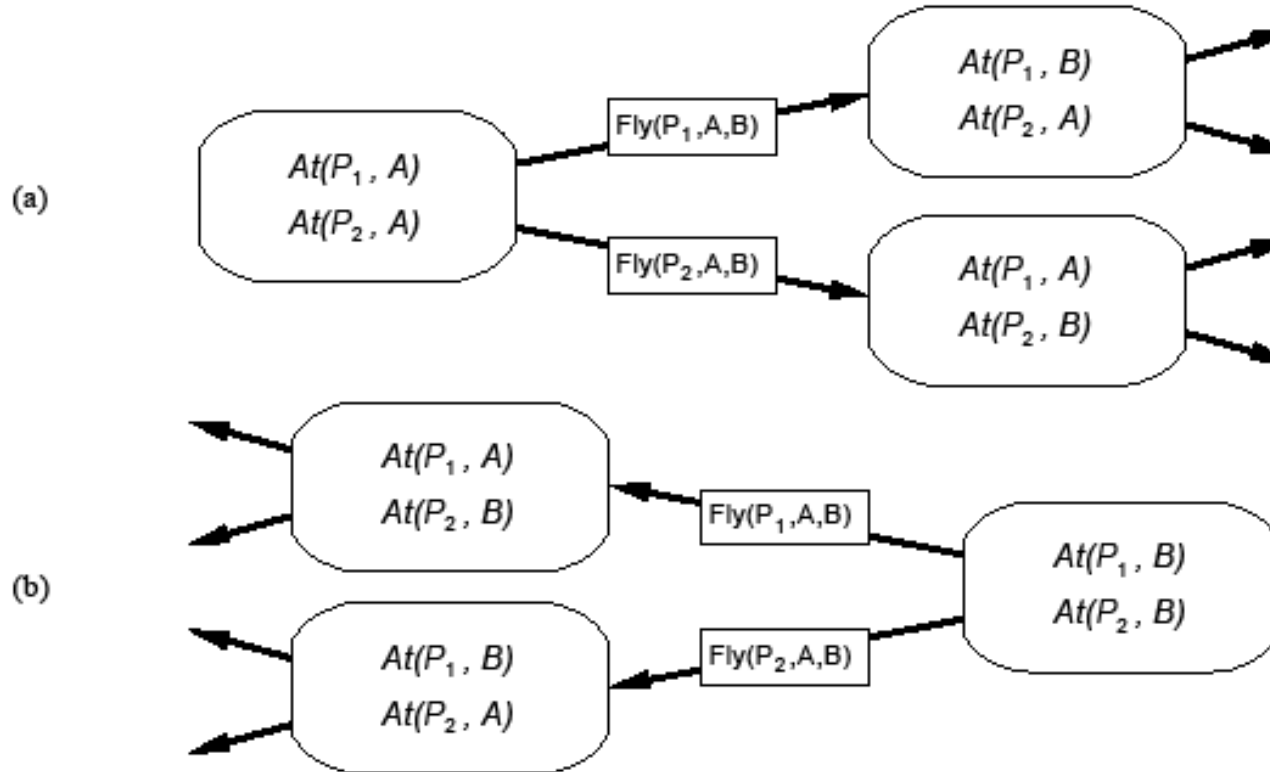
Outline

- Background
 - Situation Calculus
 - Frame, qualification, & ramification problems
- Representation language
- **Planning Algorithms**
 - **State-Space Search**
 - Partial-Order Planning (POP)
 - Planning Graphs (GRAPHPLAN)
 - SAT Planners

State-Space Search (1)

- Search the space of states (first chapters)
 - Initial state, goal test, step cost, etc.
 - Actions are the transitions between state
- Actions are invertible (why?)
 - Move forward from the initial state: Forward State-Space Search or Progression Planning
 - Move backward from goal state: Backward State-Space Search or Regression Planning

State-Space Search (2)

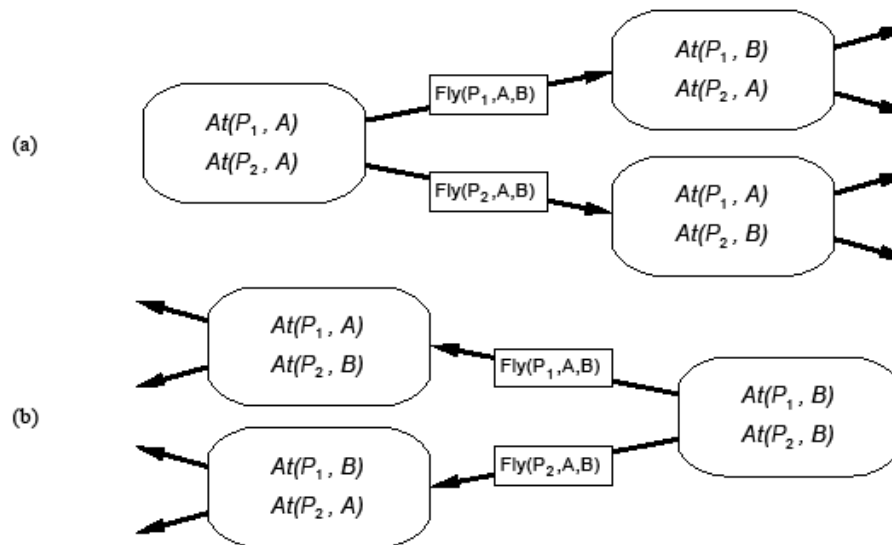


State-Space Search (3)

- Remember that the language has no functions symbols
- Thus number of states is finite
- And we can use any complete search algorithm (e.g., A^*)
 - We need an admissible heuristic
 - The solution is a path, a sequence of actions: **total-order planning**
- Problem: Space and time complexity
 - STRIPS-style planning is PSPACE-complete unless actions have
 - only positive preconditions and
 - only one literal effect

SRIPS in State-Space Search

- STRIPS representation makes it easy to focus on ‘relevant’ propositions and
 - Work backward from goal (using EFFECTS)
 - Work forward from initial state (using PRECONDITIONS)
 - Facilitating bidirectional search



Relevant Action

- An action is relevant
 - In Progression planning when its preconditions match a subset of the current state
 - In Regression planning, when its effects match a subset of the current goal state

Consistent Action

- The purpose of applying an action is to ‘achieves a desired literal’
- We should be careful that the action does not undo a desired literal (as a side effect)
- A consistent action is an action that does not undo a desired literal

Backward State-Space Search

- Given
 - A goal G description
 - An action A that is relevant and consistent
- Generate a predecessor state where
 - Positive effects (literals) of A in G are deleted
 - Precondition literals of A are added unless they already appear
 - Substituting any variables in A 's effects to match literals in G
 - Substituting any variables in A 's preconditions to match substitutions in A 's effects
- Repeat until predecessor description matches initial state

Heuristic to Speed up Search

- We can use A^* , but we need an admissible heuristic
 1. Divide-and-conquer: sub-goal independence assumption
 - Problem relaxation by removing
 2. ... all preconditions
 3. ... all preconditions and negative effects
 4. ... negative effects only: Empty-Delete-List

1. Subgoal Independence Assumption

- The cost of solving a conjunction of subgoals is the sum of the costs of solving each subgoal independently
- Optimistic
 - Where subplans interact negatively
 - Example: one action in a subplan delete goal achieved by an action in another subplan
- Pessimistic (not admissible)
 - Redundant actions in subplans can be replaced by a single action in merged plan

2. Problem Relaxation: Removing Preconditions

- Remove preconditions from action descriptions
 - All actions are applicable
 - Every literal in the goal is achievable in one step
- Number of steps to achieve the conjunction of literals in the goal is equal to the number of unsatisfied literals
- Alert
 - Some actions may achieve several literals
 - Some action may remove the effect of another action

3. Remove Preconditions & Negative Effects

- Considers only positive interactions among actions to achieve multiple subgoals
- The minimum number of actions required is the sum of the union of the actions' positive effects that satisfy the goal
- The problem is reduced to a set cover problem, which is NP-hard
 - Approximation by a greedy algorithm cannot guarantee an admissible heuristic

4. Removing Negative Effects (Only)

- Remove all negative effects of actions (no action may destroy the effects of another)
- Known as the Empty-Delete-List heuristic
- Requires running a simple planning algorithm
- Quick & effective
- Usable in progression or regression planning

Outline

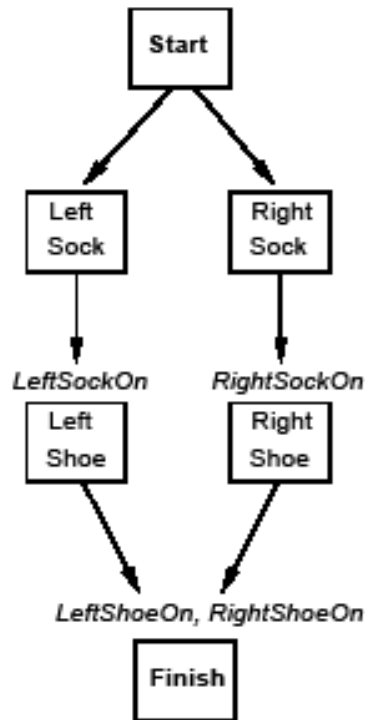
- Background
 - Situation Calculus
 - Frame, qualification, & ramification problems
- Representation language
- **Planning Algorithms**
 - State-Space Search
 - **Partial-Order Planning (POP)**
 - Planning Graphs (GRAPHPLAN)
 - SAT Planners

Partial Order Planning (POP)

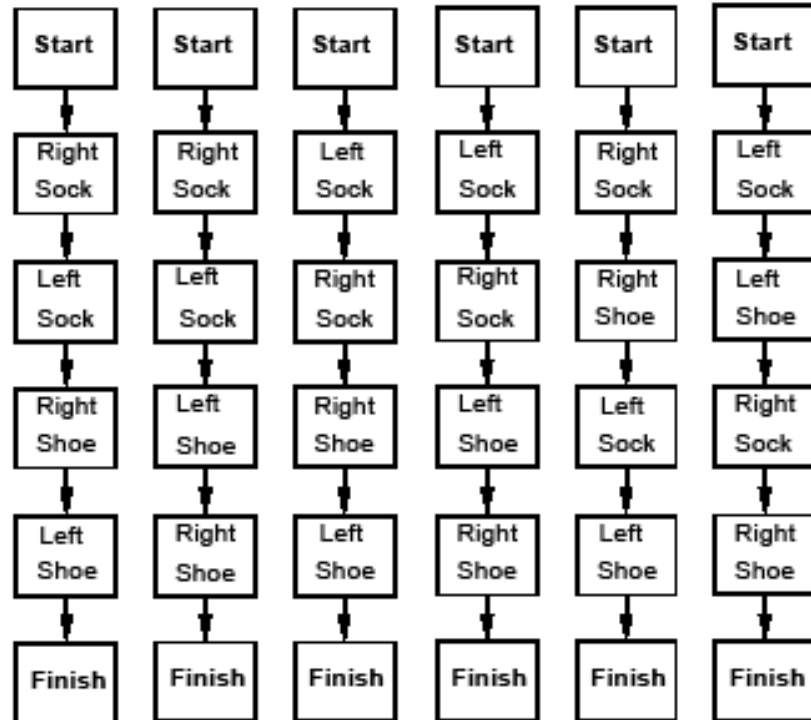
- State-space search
 - Yields totally ordered plans (linear plans)
- POP
 - Works on subproblems independently, then combines subplans
 - Example
 - Goal(RightShoeOn \wedge LeftShoeOn)
 - Init()
 - Action(RightShoe, PRECOND: RightSockOn, EFFECT: RightShoeOn)
 - Action(RightSock, EFFECT: RightSockOn)
 - Action(LeftShoe, PRECOND: LeftSockOn, EFFECT: LeftShoeOn)
 - Action(LeftSock, EFFECT: LeftSockOn)

POP Example & its linearization

Partial Order Plan:



Total Order Plans:



Components of a Plan

1. A set of **actions**
2. A set of **ordering constraints**
 - $A \prec B$ reads “A before B” but not necessarily immediately before B
 - Alert: caution to cycles $A \prec B$ and $B \prec A$
3. A set of **causal links** (protection intervals) between actions
 - $A \xrightarrow{p} B$ reads “A achieves p for B” and p must remain true from the time A is applied to the time B is applied
 - Example “RightSock $\xrightarrow{\text{RightSockOn}}$ RightShoe”
4. A set of **open preconditions**
 - Planners work to reduce the set of open preconditions to the empty set w/o introducing contradictions

Consistent Plan (POP)

- Consistent plan is a plan that has
 - No cycle in the ordering constraints
 - No conflicts with the causal links
- Solution
 - Is a consistent plan with no open preconditions
- To solve a conflict between a causal link $A \xrightarrow{p} B$ and an action C (that **clobbers**, threatens the causal link), we force C to occur outside the “protection interval” by adding
 - the constraint $C \prec A$ (**demoting** C) or
 - the constraint $B \prec C$ (**promoting** C)

Setting up the PoP

- Add dummy states

- Start

- Has no preconditions
 - Its effects are the literals of the initial state

- Finish

- Its preconditions are the literals of the goal state
 - Has no effects

- Initial Plan:

- Actions: {Start, Finish}

- Ordering constraints: {Start \prec Finish}

- Causal links: {}

- Open Preconditions: {LeftShoeOn, RightShoeOn}

Start

Literal_a, Literal_b, ...

Literal₁, Literal₂, ...

Finish

Start

LeftShoeOn, RightShoeOn

Finish

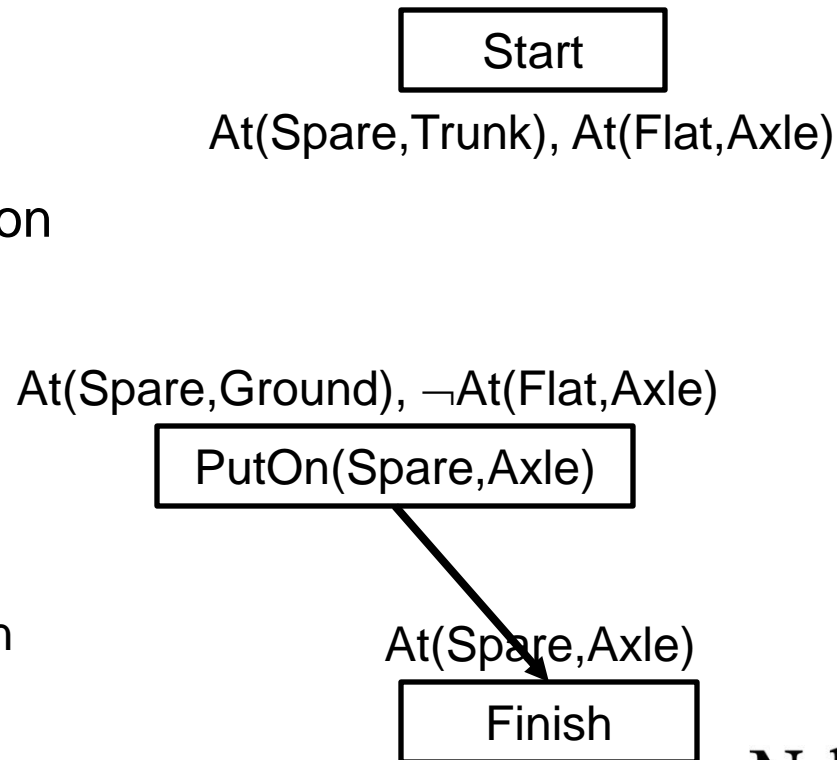
POP as a Search Problem

- The successor function arbitrarily picks one open precondition p on an action B
- For every possible consistent action A that achieves p
 - It generates a successor plan adding the causal link $A \xrightarrow{p} B$ and the ordering constraint $A \prec B$
 - If A was not in the plan, it adds $\text{Start} \prec A$ and $A \prec \text{Finish}$
 - It resolves all conflicts between
 - the new causal link and all existing actions
 - between A and all existing causal links
 - Then it adds the successor states for combination of resolved conflicts
- It repeats until no open precondition exists

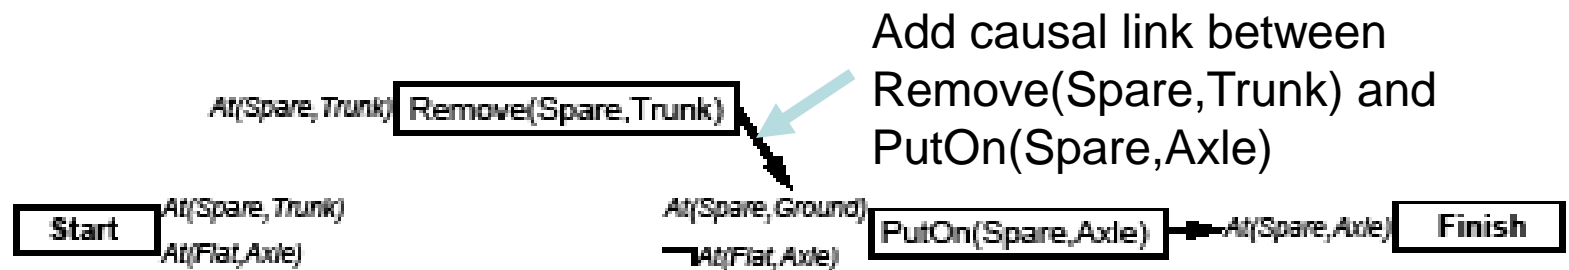
Example of POP: Flat tire problem

- See problem description in Fig 11.7 page 391


- Only one open precondition
- Only 1 applicable action



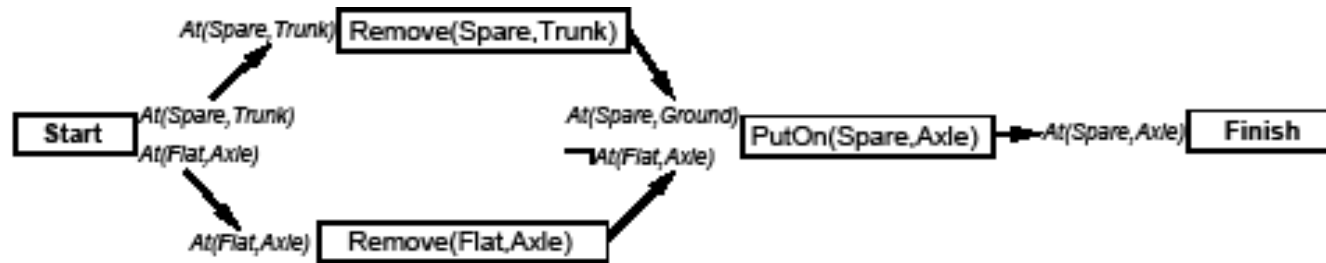
- Pick up $\text{At}(\text{Spare}, \text{Ground})$
- Choose only applicable action
 $\text{Remove}(\text{Spare}, \text{Trunk})$



- Pick up $\neg At(Flat, Axle)$
- There are 2 applicable actions: LeaveOvernight and Remove(Flat, Axle)
- Choose LeaveOvernight



- LeaveOvernight has effect $\neg At(Spare, Ground)$, which conflicts with the causal link
- We remove the conflict by forcing LeaveOvernight to occur before Remove(Spare, Trunk)
- Conflicts with effects of Remove(Spare, Trunk)
- The only way to resolve the conflict is to undo LeaveOvernight use the action Remove(Flat, Axle)



- This time, we choose Remove(Flat,Axle)
- Pick up At(Spare,Trunk) and choose Start to achieve it
- Pick up At(Flat,Axle) and choose Start to achieve it.
- We now have a complete consistent partially ordered plan

POP Algorithm (1)

- Backtrack when fails to resolve a threat or find an operator
- Causal links
 - Recognize when to abandon a doomed plan without wasting time expanding irrelevant part of the plan
 - allow early pruning of inconsistent combination of actions
- When actions include variables, we need to find appropriate substitutions
 - Typically we try to delay commitments to instantiating a variable until we have no other choice (least commitment)
- POP is sound, complete, and systematic (no repetition)

POP Algorithm (2)

- Decomposes the problem (advantage)
- But does not represent states explicitly: it is hard to design heuristic to estimate distance from goal
 - Example: Number of open preconditions – those that match the effects of the start node. Not perfect (same problems as before)
- A heuristic can be used to choose which plan to refine (which precondition to pick-up):
 - Choose the most-constrained precondition, the one satisfied by the least number of actions. Like in CSPs!
 - When no action satisfies a precondition, backtrack!
 - When only one action satisfies a precondition, pick up the precondition.

Outline

- Background
 - Situation Calculus
 - Frame, qualification, & ramification problems
- Representation language
- **Planning Algorithms**
 - State-Space Search
 - Partial-Order Planning (POP)
 - **Planning Graphs (GRAPHPLAN)**
 - SAT Planners

Planning Graph

- Is special data structure used for
 1. Deriving better heuristic estimates
 2. Extract a solution to the planning problem: GRAPHPLAN algorithm
- Is a sequence $\langle S_0, A_0, S_1, A_1, \dots, S_i \rangle$ of levels
 - Alternating state levels & action levels
 - Levels correspond to time stamps
 - Starting at initial state
 - State level is a set of (propositional) literals
 - All those literals that could be true at that level
 - Action level is a set of (propositionalized) actions
 - All those actions whose preconditions appear in the state level (ignoring all negative interactions, etc.)
- Propositionalization may yield combinatorial explosion in the presence of a large number of objects

Focus

- Building the Planning Graph
- Using it for Heuristic Estimation
- Using it for generating the plan

Example of a Planning Graph (1)

Init(Have(Cake))

Goal(Have(Cake) \wedge Eaten(Cake))

Action(Eat(Cake))

Precond: Have(Cake)

Effect: \neg Have(Cake) \wedge Eaten(Cake))

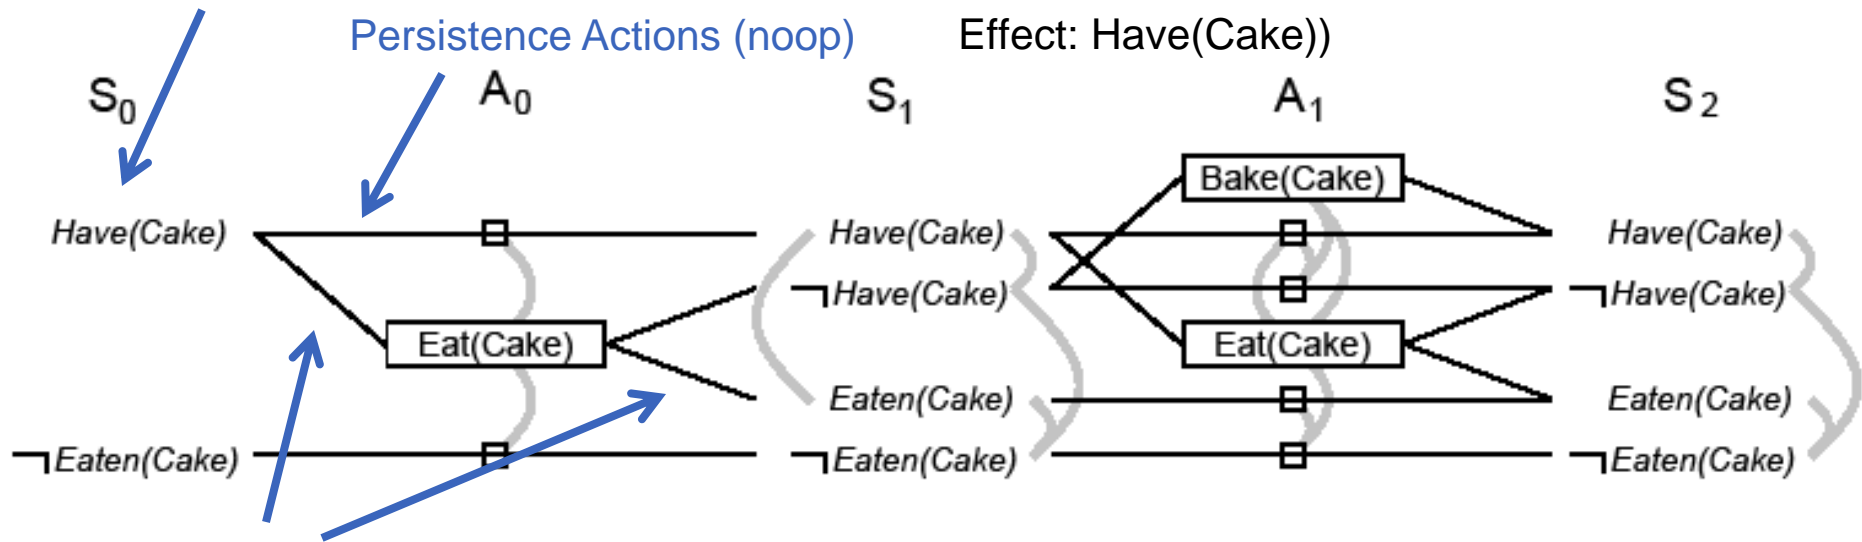
Action(Bake(Cake))

Precond: \neg Have(Cake)

Effect: Have(Cake))

Propositions true
at the initial state


Persistence Actions (noop)



Action is connected to its
preconds & effects


Example of a Planning Graph (2)

- At each state level, list all literals that may hold at that level
- At each action level, list all noops & all actions whose preconditions may hold at previous levels
- Repeat until plan 'levels off,' no new literals appears ($S_i = S_{i+1}$)
- Building the Planning Graph is a polynomial process
- Add (binary) mutual exclusion (mutex) links between conflicting actions and between conflicting literals




Mutex Links between Actions

1. **Inconsistent effects:** one action negates an effect of another
 - Eat(Cake) & noop of Have(Cake) disagree on effect Have(Cake)
2. **Interference:** An action effect negates the precondition of another
 - Eat(Cake) negates precondition of the noop of Have(Cake):
3. **Competing needs:** A precondition on an action is mutex with the precondition of another
 - Bake(Cake) & Eat(Cake): compete on Have(Cake) precondition



Mutex Links between Literals

1. Two literals are negation of each other
2. **Inconsistent support:** Each pair of actions that can achieve the two literals is mutex. Examples:
 - In S1, Have(Cake) & Eaten(Cake) are mutex
 - In S2, they are not because Bake(Cake) & the noop of Eaten(Cake) are not mutex



Focus

- Building the Planning Graph
- **Using it for Heuristic Estimation**
 - Planning graph as a relaxation of original problem
 - Easy to build (compute)
- Using it for generating the plan

Planning Graph for Heuristic Estimation

- A literal that does not appear in the final level cannot be achieved by any plan
 - State-space search: Any state containing an unachievable literal has cost $h(n)=\infty$
 - POP: Any plan with an unachievable open condition has cost $h(n)=\infty$
- The estimate cost of any goal literal is the first level at which it appears
 - Estimate is admissible for individual literals
 - Estimate can be improved by serializing the graph (serial planning graph: one action per level) by adding mutex between all actions in a given level
- The estimate of a conjunction of goal literals
 - Three heuristics: max level, level sum, set level

Estimate of Conjunction of Goal Literals

- Max-level
 - The largest level of a literal in the conjunction
 - Admissible, not very accurate
- Level sum
 - Under subgoal independence assumption, sums the level costs of the literals
 - Inadmissible, works well for largely decomposable problems
- Set level
 - Finds the level at which all literals appear w/o any pair of them being mutex
 - Dominates max-level, works extremely well on problems where there is a great deal of interaction among subplans

Focus

- Building the Planning Graph
- Using it for Heuristic Estimation
- **Using it for generating the plan**
 - GraphPlan algorithm [Blum & Furst, 95]

GRAPHPLAN algorithm

GRAPHPLAN(*problem*) **returns** *solution* or *failure*

graph \leftarrow INITIALPLANNINGGRAPH(*problem*)

goals \leftarrow GOALS[*problem*]

loop do

if *goals* all non-mutex in last level of *graph* **then do**

solution \leftarrow EXTRACTSOLUTION(*graph*,*goals*,LENGTH(*graph*))

if *solution* \neq *failure* **then return** *solution*

else if NOSOLUTIONPOSSIBLE(*graph*) **then return** *failure*

graph \leftarrow EXPANDGRAPH (*graph*,*problem*)

- Two main stages
 1. Extract solution
 2. Expand the graph

Example of GRAPHPLAN Execution (1)

- $\text{At}(\text{Spare}, \text{Axle})$ is not in S_0
- No need to extract solution
- Expand the plan

S_0
 $\text{At}(\text{Spare}, \text{Trunk})$

$\text{At}(\text{Flat}, \text{Axle})$


$\neg \text{At}(\text{Spare}, \text{Axle})$

$\neg \text{At}(\text{Flat}, \text{Ground})$

$\neg \text{At}(\text{Spare}, \text{Ground})$


Example of GRAPHPLAN Execution (2)

- Three actions are applicable
- 3 actions and 5 noops are added
- Mutex links are added
- $At(Spare, Axle)$ still not in S_1
- Plan is expanded



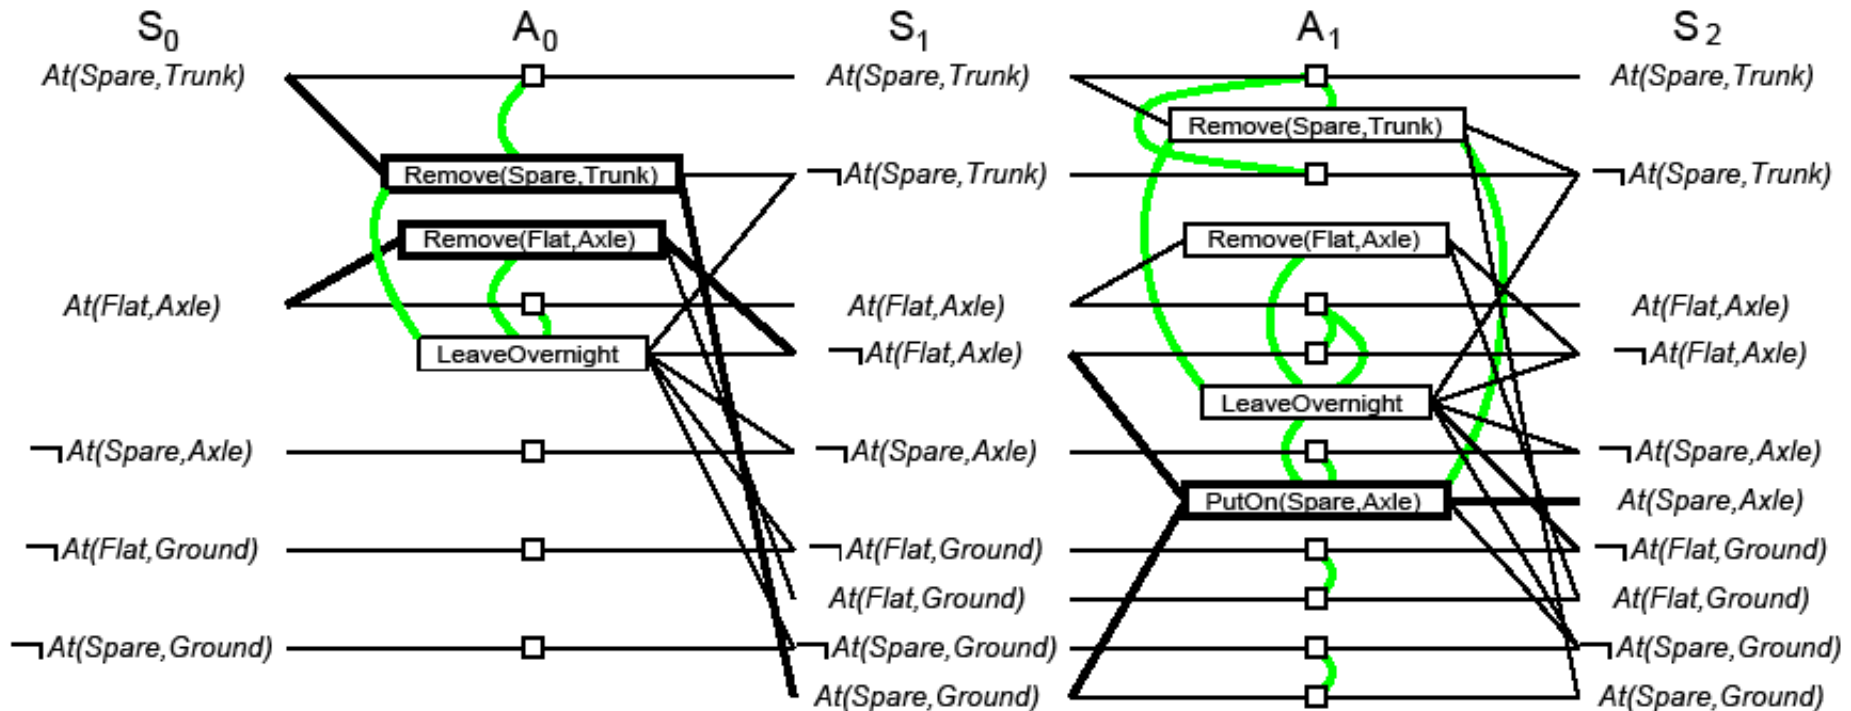
Example of GRAPHPLAN Execution (3)

- Illustrates well mutex links: inconsistent effects, interference, competing needs, inconsistent support



Solution Extraction (Backward)

1. Solve a Boolean CSP: Variables are actions, domains are {0=out of plan, 1=in plan}, constraints are mutex
2. Search problem from last level backward



Backtrack Search for Solution Extraction

- Starting at the highest fact level
 - Each goal is put in a goal list for the current fact layer
 - Search iterates thru each fact in the goal list trying to find an action to support it which is not mutex with any other chosen action
 - When an action is chosen, its preconditions are added to the goal list of the lower level
 - When all facts in the goal list of the current level have a consistent assignment of actions, the search moves to the next level
- Search backtracks to the previous level when it fails to assign an action to each fact in the goal list at a given level
- Search succeeds when the first level is reached.

Termination of GRAPHPLAN

- GRAPHPLAN is guaranteed to terminate
 - Literal increase monotonically
 - Actions increase monotonically
 - Mutexes decrease monotonically
- A solution is guaranteed not to exist when
 - The graph levels off with all goals present & non-mutex, and
 - EXTRACTSOLUTION fails to find solution

Optimality of GRAPHPLAN

- The plans generated by GRAPHPLAN
 - Are optimal in the number of steps needed to execute the plan
 - Not necessarily optimal in the number of actions in the plan (GRAPHPLAN produces partially ordered plans)

Outline

- Background
 - Situation Calculus
 - Frame, qualification, & ramification problems
- Representation language
- Planning Algorithms
 - State-Space Search
 - Partial-Order Planning (POP)
 - Planning Graphs (GRAPHPLAN)
 - SAT Planners